

How to Read and Summarize Existing Research Studies?

Tsz Nam Chan

edisonchan@szu.edu.cn

College of Computer Science & Software Engineering

Why Do We Need to Summarize Existing Research Studies?

- You are not familiar with that topic.
 - You are interested in conducting research for it.
 - You want to understand the state-of-the-art solutions.
- You want to write a survey paper for that topic.
- You want to conduct a tutorial (write a tutorial paper) for that topic.
- You need to write the “Related Work” section.
- Your PhD supervisor asks you to do so.

Steps for Summarizing Research Papers

1. Read papers.
2. Categorize those papers into different groups.
3. Summarize different groups of research papers.
 1. Core ideas of different groups.
 2. The main differences (e.g., advantages and disadvantages) between different groups of papers.

How to Read Papers?

- Looks easy. (But it is very difficult.)
- Not necessary (and not feasible) to fully read every paper.


- Inge Vejsbjerg, Elizabeth M. Daly, Rahul Nair, Svetoslav Nizhnichenkov:
Interactive Human-Centric Bias Mitigation. 23838-23840
- Jiaying Wang, Shuailing Hao, Jing Shan, Xiaoxu Song:
Visual Language - Let the Product Say What You Want. 23841-23843
- Kuang-Da Wang, Yu-Tse Chen, Yu-Heng Lin, Wei-Yao Wang, Wen-Chih Peng:
The CoachAI Badminton Environment: Bridging the Gap between a Reinforcement Learning Environment and Real-World Badminton Games. 23844-23846
- Umer Waqas, Yunwan Jeon, Donghun Lee:
Virtual Try-On: Real-Time Interactive Hybrid Network with High-Fidelity. 23847-23849
- Lianlong Wu, Seewon Choi, Daniel Raggi, Aaron Stockdill, Grecia Garcia Garcia, Fiorenzo Colarusso, Peter C.-H. Cheng, Mateja Jamnik:
Generation of Visual Representations for Multi-Modal Mathematical Knowledge. 23850-23852
- Tiancheng Zhang, Shaoyuan Huang, Cheng Zhang, Xiaofei Wang, Wenyu Wang:
EasyTS: The Express Lane to Long Time Series Forecasting. 23853-23855
- Zeyuan Zhang, Tanmay Laud, Zihang He, Xiaojie Chen, Xinshuang Liu, Zhouhang Xie, Julian J. McAuley, Zhankui He:
RecWizard: A Toolkit for Conversational Recommendation with Modular, Portable Models and Interactive User Interface. 23856-23858
- Runcong Zhao, Wenjia Zhang, Jiazheng Li, Lixing Zhu, Yanran Li, Yulan He, Lin Gui:
NarrativePlay: An Automated System for Crafting Visual Worlds in Novels for Role-Playing. 23859-23861

In AAAI 2024

- Fully reading a paper does not indicate that you understand that paper.

How to Read Papers?

- Read title.
 - You can filter a large portion of papers based on this.
- Read Abstract and Introduction.
 - You can categorize those papers into different groups.
 - You can identify important papers.
- Read Preliminaries and Related Work.
 - You can (further) categorize those papers into different groups.
 - You can identify some missing papers.
- Only need to fully read a few (important) papers.



How to Categorize Papers into Different Groups?

- Understanding the main technical novelty of this paper.
 - This can normally be found in the abstract and introduction sections.
- What are the main novelties of this paper?
 - Develop a new sampling method.
 - Achieve non-trivial sampling guarantees.
- Which category does this paper belong to?
 - Sampling

Improved Coresets for Kernel Density Estimates

Jeff M. Phillips*
University of Utah

Wai Ming Tai
University of Utah

Abstract

We study the construction of coresets for kernel density estimates. That is we show how to approximate the kernel density estimate described by a large point set with another kernel density estimate with a much smaller point set. For characteristic kernels (including Gaussian and Laplace kernels), our approximation preserves the L_∞ error between kernel density estimates within error ε , with coreset size $4/\varepsilon^2$, but no other aspects of the data, including the dimension, the diameter of the point set, or the bandwidth of the kernel common to other approximations. When the dimension is unrestricted, we show this bound is tight for these kernels as well as a much broader set.

This work provides a careful analysis of the iterative Frank-Wolfe algorithm adapted to this context, an algorithm called *kernel herding*. This analysis unites a broad line of work that spans statistics, machine learning, and geometry.

When the dimension d is constant, we demonstrate much tighter bounds on the size of the coreset specifically for Gaussian kernels, showing that it is bounded by the size of the coreset for axis-aligned rectangles. Currently the best known constructive bound is $O(\frac{1}{\varepsilon} \log^d \frac{1}{\varepsilon})$, and non-constructively, this can be improved by $\sqrt{\log \frac{1}{\varepsilon}}$. This improves the best constant dimension bounds polynomially for $d > 3$.

infinite dimensional function spaces (each KDE_P is a point in such a space). From these techniques grew much of non-linear data analysis (e.g., kernel PCA, kernel SVM). In particular, an object in the RKHS called the *kernel mean* is another representation of KDE_P , and its sparse approximation plays a critical role in distribution hypothesis testing [15, 16], Markov random fields [4], and even political data analysis [32]. Through a simple argument (described below), the standard approximation of the kernel mean in the RKHS implies a L_∞ approximation bound of the kernel density estimate in \mathbb{R}^d [4, 34] (which is stronger than the L_1 and L_2 variants [37]).

More recently, the sparse approximation of a kernel density estimate has gained interest from the computational geometry community for its connections in topological data analysis [29, 9], coresets [27], and discrepancy theory [17].

In this paper, we provide strong connections between all of these storylines, and in particular provide a simpler analysis of the common sparse kernel mean approximation techniques with application to the strong L_∞ -error coresets of kernel density estimates. With unrestricted dimensions, we show our bounds for KDEs are tight, and in constant dimensions of at least 3, we polynomially improve the best known bounds so they are now tight up to poly-log factors.

How to Categorize Papers into Different Groups?

- What are the main novelties of this paper?
 - Develop a new sampling method.
 - Achieve non-trivial sampling guarantees.
- Which category does this paper belong to?
 - Sampling

1 INTRODUCTION

DATA is collected at ever-increasing sizes, and for many datasets, each data point has geo-spatial locations (e.g., either (x, y) -coordinates, or latitudes and longitudes). Examples include population tracking data, geo-located social media contributions, seismic data, crime data, and weather station data. The availability of such detailed datasets enables analysts to ask more complex and specific questions. These have applications in wide ranging areas including biosurveillance, epidemiology, economics, ecology environmental management, public policy and safety, transportation design and monitoring, geology, and climatology. Truly large datasets, however, cannot be simply plotted, since they typically exceed the number of pixels available for plotting, the available storage space, and/or the available bandwidth necessary to transfer the data.

A common way to manage and visualize such large, complex spatial data is to represent it using a continuous, smoothed function, typically a kernel density estimate [1], [2] (KDE). A KDE is a statistically and spatially robust method to represent a continuous density using only a discrete set of sample points. Informally, this can be thought of as a continuous average over all choices of histograms, which avoid some instability issues that arise in histograms due to discretization boundaries. Or it is a convolution of all data points with a continuous smoothing function. For a formal definition, we first require a kernel $K : \mathbb{R}^2 \times \mathbb{R}^2 \rightarrow \mathbb{R}$; we will use the Gaussian kernel $K(p, x) = e^{-\|p-x\|^2}$, the most

-
- Y. Zheng is with Visa Research, Palo Alto, CA 94306 USA.
E-mail: yanzh.cs@gmail.com.
 - Y. Ou is with Expedia, Inc., Bellevue, WA 98004 USA.
E-mail: olly93219@outlook.com.
 - A. Lex and J. Phillips are with the University of Utah, Salt Lake City, UT 84112 USA. E-mail: alex@sci.utah.edu, jeffp@cs.utah.edu.

Manuscript received 20 Dec. 2017; revised 16 Sept. 2018; accepted 21 Mar. 2019. Date of publication 29 Apr. 2019; date of current version 29 June 2021. (Corresponding author: Yan Zheng).

Recommended for acceptance by A. Perer and C. Scheidegger.
Digital Object Identifier no. 10.1109/TBDB.2019.2913655

common and pervasive kernel. Then, given a planar point set $P \subset \mathbb{R}^2$, the kernel density estimate is defined at any query point $x \in \mathbb{R}^2$ as

$$\text{KDE}_P(x) = \frac{1}{|P|} \sum_{p \in P} K(p, x).$$

This allows regions with more points nearby (i.e., points x with a large value $K(p, x)$ for many p in P) to have a large density value, and this function is smooth and in general nicely behaved in many contexts. Using this function summarizes the data, and avoids the over-plotting and obfuscation issues demonstrated in Fig. 1 (left). However, just computing $\text{KDE}_P(x)$ for a single value x requires $O(|P|)$ time; it iterates over all data points summing their contributions. While these values can be precomputed and mapped to a bitmap, visually interacting with a KDE e.g., to query and filter, would then require expensive reaggregating. For instance, as a user zooms in on a region of interest, ideally the visual interface should increase the resolution, and possibly shift the grid boundaries. This would require recomputing each of these visible pixel values in $O(|P|)$ time each.

Towards alleviating these issues, we propose to use coresets for KDEs. In general, a coreset Q is a proxy for a large set P ; it is a carefully designed small subset of a very large dataset P where Q retains properties from P as accurately as possible. In particular, in many cases the size of Q depends only on a desired minimum level of accuracy, not the size of the original dataset P . This implies that even if the full dataset grows, the size of the coreset required to represent a phenomenon stays fixed. This also holds when P represents a continuous quantity (like the locus of points along a road network, or a spread of particulates from a forest fire) and Q constitutes some carefully placed representative points [3]. Fig. 1 shows a dataset P with 700 thousand points and its coreset from all reported crimes in Philadelphia from 2005–2014. For more details on variations and constructions, refer to recent surveys [4], [5].

How to Categorize Papers into Different Groups?

- What are the main novelties of this paper?
 - Develop a new function approximation method.
 - Achieve non-trivial approximation guarantees.
- Which category does this paper belong to?
 - Function approximation




Figure 1: A color map for motor vehicle thefts (black dots) in Arlington, Texas in 2007 (Cropped from [23])

$\mathcal{F}_P(q)$ [46]. Equation 1 shows one example of $\mathcal{F}_P(q)$ with Gaussian kernel, where P and $dist(q, p_i)$ are the set of two-dimensional data points and Euclidean distance respectively.

$$\mathcal{F}_P(q) = \sum_{p_i \in P} w_i \cdot \exp(-\gamma dist(q, p_i)^2) \quad (1)$$

In this paper, we will also consider $\mathcal{F}_P(q)$ with other kernel functions in Section 5. As a remark, all kernel functions, that we consider in this paper, are adopted in famous software, e.g., Scikit-learn [38] and QGIS [43].

A higher $\mathcal{F}_P(q)$ value indicates a higher density of data points in the region around q . The above KDE function is computationally expensive to compute. Given a data set with 1 million 2D points, KDV involves over 2 trillion operations [39] on a 1920×1080 screen. As pointed out in [16, 19, 55, 58, 59], KDV cannot scale well to handle many data points and display of color maps on high-resolution screens. To address this problem, researchers have proposed two variants of KDV, which aim to improve its performance:

• ϵ KDV: This is an approximate version of KDV. A relative error parameter, ϵ , is used, such that for each pixel q , the pixel color is within $(1 \pm \epsilon)$ of $\mathcal{F}_P(q)$. Figure 2a shows a color map generated by the original (exact) KDV, while Figure 2b illustrates the corresponding color map for ϵ KDV with ϵ equal to 0.01. As we can see, the two color maps do not look different. ϵ KDV runs faster than exact KDV [10, 20, 57–59], and is also supported in data analytics software (e.g., Scikit-learn [38]).

• τ KDV: In tasks such as hotspot detection [6, 23], a data visualization user only needs to know which spatial region has a high density (i.e., hotspot), and not the other areas. One such hotspot is the red region in Figure 1. A color map with two colors are already sufficient. Figure 2c shows such a color map. To generate this color map, the τ KDV can be used, where a threshold, τ , determines the color of a pixel: a color for q when $\mathcal{F}_P(q) \geq \tau$ (to indicate high density), and another color otherwise. This method, recently studied in [10, 16], is shown to be faster than exact KDV.

Although ϵ KDV and τ KDV perform better than exact KDV, they still require a lot of time. On a $270k$ -point crime dataset [2], displaying a color map on a screen with 1280×960 pixels takes over an hour for most methods, including the ϵ KDV solution implemented in Scikit-learn. In fact, these existing methods often cannot deliver *real-time* performance, which allows color maps to be generated quickly, thereby saving the precious waiting time of data analysts.

Our contributions. In this paper, we develop a solution, called QUAD, in order to improve the performance of ϵ KDV and τ KDV. The main idea is to derive lower and upper bounds of the KDE function (i.e., Equation 1) in terms of quadratic functions (cf. Section 4). These *quadratic bounds* are theoretically tighter than the existing ones (in ϵ KDE [20], ϵ KDC [16], and KARL [10]), enabling faster pruning. In addition, many KDV-based applications [14, 18, 23, 30] also utilize other kernel functions, including triangular, cosine kernels etc. Therefore, we extend our techniques to support other kernel functions (cf. Section 5), which cannot be supported by the state-of-the-art solution, KARL [10].

In our experiments on large datasets in a single machine, QUAD is at least one-order-of-magnitude faster than existing solutions. For ϵ KDV, QUAD takes 100–1000 sec to generate color map for each large-scale dataset ($0.17M$ to $7M$) with 2560×1920 pixels, using small relative error $\epsilon = 0.01$. However, most of the other methods fail to generate the color map within 2 hours under the same setting. For τ KDV, QUAD can achieve nearly 10 sec with 1280×960 pixels, using different thresholds.

We further adopt a progressive visualization framework for KDV (cf. Section 6), in order to continuously output partial visualization results (by increasing the resolution). A user can terminate the process anytime, once the partial visualization results are satisfactory, instead of waiting for the precise color map to be generated. Experiment results show that we can achieve real-time (0.5 sec) in single machine without using GPU and parallel computation by combining this framework with our solution QUAD.

The rest of the paper is organized as follows. We first review existing work in Section 2. We then discuss the background in Section 3. Later, we present quadratic bound functions for KDE in Section 4. After that, we extend our quadratic bounds to other kernel functions in Section 5. We then discuss our progressive visualization framework for KDV in Section 6. Lastly, we show our results in Section 7, and conclude in Section 8. All the proofs are shown in Section 9.

2 RELATED WORK

Kernel density visualization (KDV) is widely used in many application domains, such as: ecological modeling [32, 33], crime [6, 23, 56] or traffic hotspot detection [48, 52, 54], chemical geology [49] and physical modeling [13]. For each application, they either need to compute the approximate kernel density values with theoretical guarantee [20] (ϵ KDV) or

How to Categorize Papers into Different Groups?

- What is good categorization of papers?
 - Each category should not have too many/too few papers.
 - Every paper in a category should share a lot of similar properties.
 - Different categories should have significant differences.
- Bad categorization ☹
 1. Paper 1 uses a linear function to approximate the kernel density function.
 2. Paper 2 uses a quadratic function to approximate the kernel density function.
 3. Paper 3 uses a polynomial function to approximate the kernel density function.

What's wrong with this?

How to Categorize Papers into Different Groups?

- Good categorization 😊
 1. Paper A₁, Paper A₂, Paper A₃,... adopt the function approximation method for handling this problem.
 2. Paper B₁, Paper B₂, Paper B₃,... adopt the sampling method for handling this problem.
 3. Paper C₁, Paper C₂, Paper C₃,... adopt the computational geometry method for handling this problem.

Why is this categorization good?

How to Summarize Different Groups of Research Papers?

- Write down the core ideas of each group of research papers in your own words.
- Clearly point out the advantages and disadvantages.
 - Not just a copy-and-paste work.
 - Write these differences (or advantages/disadvantages) in your own words.
- Need to clearly understand different groups of papers.

How to Summarize Different Groups of Research Papers?

- Summarization of the sampling methods from some weak researchers.

“The main contribution is that they provide strong connections between all of these storylines, and in particular provide a simpler analysis of the common sparse kernel mean approximation techniques with application to the strong L_∞ -error coresets of kernel density estimates. With unrestricted dimensions, they show their bounds for KDEs are tight, and in constant dimensions of at least 3, they polynomially improve the best known bounds so these bounds are now tight up to poly-log factors.”

No understanding

Improved Coresets for Kernel Density Estimates

Jeff M. Phillips*
University of Utah

Wai Ming Tai
University of Utah

Abstract

We study the construction of coresets for kernel density estimates. That is we show how to approximate the kernel density estimate described by a large point set with another kernel density estimate with a much smaller point set. For characteristic kernels (including Gaussian and Laplace kernels), our approximation preserves the L_∞ error between kernel density estimates within error ε , with coreset size $4/\varepsilon^2$, but no other aspects of the data, including the dimension, the diameter of the point set, or the bandwidth of the kernel common to other approximations. When the dimension is unrestricted, we show this bound is tight for these kernels as well as a much broader set.

This work provides a careful analysis of the iterative Frank-Wolfe algorithm adapted to this context, an algorithm called *kernel herding*. This analysis unites a broad line of work that spans statistics, machine learning, and geometry.

When the dimension d is constant, we demonstrate much tighter bounds on the size of the coreset specifically for Gaussian kernels, showing that it is bounded by the size of the coreset for axis-aligned rectangles. Currently the best known constructive bound is $O(\frac{1}{\varepsilon} \log^d \frac{1}{\varepsilon})$, and non-constructively, this can be improved by $\sqrt{\log \frac{1}{\varepsilon}}$. This improves the best constant dimension bounds polynomially for $d \geq 3$.

1 Introduction

A kernel density estimate [26] of a point set $P \subset \mathbb{R}^d$ smooths out the point set to create a continuous function $\text{KDE}_P : \mathbb{R}^d \rightarrow \mathbb{R}$. This object has a rich history and many applications in statistical data analysis [33, 7, 31], with many results around the question of if P is drawn iid from an unknown distribution ψ , how well can KDE_P converge to ψ as a function of $|P|$ (mainly in the L_2 [33, 31] and L_1 [7] sense).

Then kernel techniques in machine learning [30] developed the connection of kernel density estimates to reproducing kernel Hilbert spaces (RKHS), which are

infinite dimensional function spaces (each KDE_P is a point in such a space). From these techniques grew much of non-linear data analysis (e.g., kernel PCA, kernel SVM). In particular, an object in the RKHS called the *kernel mean* is another representation of KDE_P , and its sparse approximation plays a critical role in distribution hypothesis testing [15, 16], Markov random fields [4], and even political data analysis [32]. Through a simple argument (described below), the standard approximation of the kernel mean in the RKHS implies an L_∞ approximation bound of the kernel density estimate in \mathbb{R}^d [4, 34] (which is stronger than the L_1 and L_2 variants [37]).

More recently, the sparse approximation of a kernel density estimate has gained interest from the computational geometry community for its connections in topological data analysis [29, 9], coresets [27], and discrepancy theory [17].

In this paper, we provide strong connections between all of these storylines, and in particular provide a simpler analysis of the common sparse kernel mean approximation techniques with application to the strong L_∞ -error coresets of kernel density estimates. With unrestricted dimensions, we show our bounds for KDEs are tight, and in constant dimensions of at least 3, we polynomially improve the best known bounds so they are now tight up to poly-log factors.

Formal definitions. For a point set $P \subset \mathbb{R}^d$ of size n and a kernel $K : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$, a *kernel density estimate* KDE_P at $x \in \mathbb{R}^d$ is defined $\text{KDE}_P(x) = \frac{1}{|P|} \sum_{p \in P} K(x, p)$. Our goal is to construct a subset $Q \subset P$, and bound its size, so that its KDE has ε -bounded L_∞ error:

$$\|\text{KDE}_P - \text{KDE}_Q\|_\infty = \max_{x \in \mathbb{R}^d} |\text{KDE}_P(x) - \text{KDE}_Q(x)| \leq \varepsilon.$$

We call such a subset Q an *ε -coreset of a kernel range space* ((P, \mathcal{K})) (or just an *ε -kernel coreset* for short), where \mathcal{K} is the set of all functions $K(x, \cdot)$ represented by a fixed kernel K and an arbitrary center point $x \in \mathbb{R}^d$.

While there is not one standard definition of a kernel, many of these kernels have properties that unite them. Common examples are the Gaussian kernel $K(x, p) = \exp(-\|x - p\|^2/\sigma^2)$, the Laplace kernel

*Thanks to supported by NSF CCF-1350888, IIS-1251019, ACI-1443046, CNS-1514520, and CNS-1564287.


How to Summarize Different Groups of Research Papers?

- Summarization of the sampling methods from professional researchers.

Data Sampling

- Consider the kernel density function (with the Gaussian kernel).

$$\mathcal{F}_P(\mathbf{q}) = \sum_{\mathbf{p} \in P} w \cdot \exp\left(-\frac{1}{b^2} \text{dist}(\mathbf{q}, \mathbf{p})^2\right)$$



- Compute the modified kernel density function based on the sampled dataset S .

$$\mathcal{F}_S^{(M)}(\mathbf{q}) = \sum_{\mathbf{p}_i \in S} w_i \cdot \exp\left(-\frac{1}{b^2} \text{dist}(\mathbf{q}, \mathbf{p}_i)^2\right)$$

Advantages and Disadvantages of Data Sampling

- Advantages ☺
 - Can achieve probabilistic approximation guarantees for generating KDV.
 - Can reduce the worst-case time complexity for generating KDV.
 - Can handle all kernel functions.
- Disadvantages ☹
 - Cannot achieve exact solution.
 - Can still be slow for generating KDV.
 - Can degrade the practical visualization quality.

With deep understanding

How to Summarize the Main Differences Between Different Groups of Papers?

- Summarization needs your own view (the most difficult part for students).
- Be brave
 - Academic freedom protects you. ☺
 - No one will laugh at you.
 - No one will blame you even if you are wrong.
 - If they laugh at you or blame you, you can simply say this is academic freedom (I can say whatever I say as long as they are ethical or they do not break the law.).
 - Don't be afraid if it is wrong.
 - “Everything you say” is only based on the best of your knowledge at that time.
 - You only need to say “I am sorry.” or “I learn it.” when it turns out to be wrong.

Examples of Summarization of Research Papers

- “Large-scale Geospatial Analytics: Problems, Challenges, and Opportunities” in SIGMOD 2023 Tutorial (from pages 27 - 36 in the slides from this [link](#))
- “Kernel Density Visualization for Big Geospatial Data: Algorithms and Applications” in MDM 2023 Tutorial (from pages 18 - 32 in the slides from this [link](#))

More Suggestions: Reading Papers

- Not necessary to read every word of the paper.
- Don't believe everything in the paper.
- Don't destroy your confidence.

More Suggestions: Reading Papers

- Suppose that you want to improve the efficiency of Kernel Density Visualization and you see this paper. What do you think?
- Some students: This is the end of the world. 
- Some good researchers: Worth for investigating more. 

ABSTRACT

Kernel density visualization, or KDV, is used to view and understand data points in various domains, including traffic or crime hotspot detection, ecological modeling, chemical geology, and physical modeling. Existing solutions, which are based on computing kernel density (KDE) functions, are computationally expensive. Our goal is to improve the performance of KDV, in order to support large datasets (e.g., one million points) and high screen resolutions (e.g., 1280×960 pixels). We examine two widely-used variants of KDV, namely approximate kernel density visualization (ϵ KDV) and thresholded kernel density visualization (τ KDV). For these two operations, we develop fast solution, called QUAD, by deriving quadratic bounds of KDE functions for different types of kernel functions, including Gaussian, triangular etc. We further adopt a progressive visualization framework for KDV, in order to stream partial visualization results to users continuously. Extensive experiment results show that our new KDV techniques can provide at least one-order-of-magnitude speedup over existing methods, without degrading visualization quality. We further show that QUAD can produce the reasonable visualization results in real-time (0.5 sec) by combining the progressive visualization framework in single machine setting without using GPU and parallel computation.

More Suggestions: Reading Papers

- They focus on developing approximate algorithms. How about exact algorithms?
- Maybe they use some “good” datasets.
- It is possible that progressive visualization framework provides bad results.


DOUBT

ABSTRACT

Kernel density visualization, or KDV, is used to view and understand data points in various domains, including traffic or crime hotspot detection, ecological modeling, chemical geology, and physical modeling. Existing solutions, which are based on computing kernel density (KDE) functions, are computationally expensive. Our goal is to improve the performance of KDV, in order to support large datasets (e.g., one million points) and high screen resolutions (e.g., 1280×960 pixels). We examine two widely-used variants of KDV, namely approximate kernel density visualization (ϵ KDV) and thresholded kernel density visualization (τ KDV). For these two operations, we develop fast solution, called QUAD, by deriving quadratic bounds of KDE functions for different types of kernel functions, including Gaussian, triangular etc. We further adopt a progressive visualization framework for KDV, in order to stream partial visualization results to users continuously. Extensive experiment results show that our new KDV techniques can provide at least one-order-of-magnitude speedup over existing methods, without degrading visualization quality. We further show that QUAD can produce the reasonable visualization results in real-time (0.5 sec) by combining the progressive visualization framework in single machine setting without using GPU and parallel computation.

More Suggestions: Reading Papers

- Only focus on generating KDV in the planar space.
- How about the road network space?
- How about the temporal part of those data points?



More Suggestions: Reading Papers

- Only focus on Gaussian kernel function and those kernel functions in the following table.

| Kernel function | Equation ($\mathcal{K}(q, p)$) | Used in |
|-----------------|--|--------------|
| Triangular | $\max(1 - \gamma \cdot dist(q, p), 0)$ | [18, 23] |
| Cosine | $\begin{cases} \cos(\gamma dist(q, p)) & \text{if } dist(q, p) \leq \frac{\pi}{2\gamma} \\ 0 & \text{otherwise} \end{cases}$ | [14, 23, 30] |
| Exponential | $\exp(-\gamma \cdot dist(q, p))$ | [23] |

- How about other kernels? (Epanechnikov kernel and quartic kernel are also very famous.)

NOT NECESSARY TO FOLLOW THE SAME SETTING

How to Find Other Settings?

- Keep answering questions when you are reading papers.
 - How about the road network space? [a]
 - How about the temporal part of those data points? [b]
 - How about other kernels? [a, b, c]


[a] **Tsz Nam Chan**, Zhe Li, Leong Hou U, Jianliang Xu, Reynold Cheng: “Fast Augmentation Algorithms for Network Kernel Density Visualization” **PVLDB** 2021 (vol 14), pages 1503-1516.

[b] **Tsz Nam Chan**, Pak Lon Ip, Leong Hou U, Byron Choi, Jianliang Xu: “SWS: A Complexity-Optimized Solution for Spatial-Temporal Kernel Density Visualization” **PVLDB** 2022 (vol 15), pages 814-827.


[c] **Tsz Nam Chan**, Leong Hou U, Byron Choi, Jianliang Xu: “SLAM: Efficient Sweep Line Algorithms for Kernel Density Visualization” **SIGMOD** 2022, pages 2120-2134.

- Possible for you to have new problem settings (i.e., new papers) by answering those questions.

More Suggestions: Writing Papers




Weak researcher: There will be no topic to work with when I get a tenure-track position. I need to retain some of them for publishing when I get the tenure-track position.




Strong researcher: The more I write, the more topics I will have.

More Suggestions: Reading and Writing Papers




Junior students



"Senior" students


More Suggestions: Reading and Writing Papers

1. Want to work on KDV.
2. Start reading some papers in KDV.
3. I have learnt somethings. Write it down in the draft. Moreover, there are still some questions (e.g., Can indexing structures be used in KDV?).
4. Check the literature again and figure out the answers and ask some additional questions...




More Suggestions: Reading and Writing Papers

5. Are those indexing structure good?
(Do some research)
6. I have learnt somethings. Write it down in the draft. Moreover, there are still some other questions.
7. Check the literature again and figure out the answers and ask some additional questions...
8. Come up with new idea.
9. Write it down in the draft and do research.



Our New Book



Tsz Nam Chan, Dingming Wu

Mastering the Academic Writing
Mindset

A Guide to Crafting Computer Science Papers

September 12, 2025

Springer Nature

Our New Book

- Hard copy can be bought in some bookstores (e.g., Amazon).
 - My signature is free. ☺
- Soft copy is available online (open access).
- Available on 31st March 2026.

